

Homework 3: ELF files

This assignment will make you more familiar with organisation of ELF files. Technically, you can do this assignment on any operating system that supports the Unix API (Linux CADE machines, your laptop that runs Linux or Linux VM, and even MacOS, etc.). **You don't need to set up xv6 for this assignment** Submit your programs and the shell through Gradescope (see instructions at the bottom of this page).

NOTE: YOU CANNOT PUBLICLY RELEASE SOLUTIONS TO THIS HOMEWORK. It's ok to show your work to your future employer as a private Git repo, however any public release is prohibited. For **Mac / OSX** users. The support of 32 bit applications is deprecated in the latest version of your system. So if you already updated your system to MacOS Catalina or have updated your XCode then we recommend you to do the homework at the CADE machines.

Part 1: Take a look at ELF files

Download the [main.c](#), and [elf.c](#), and look over them. At a high level this homework asks you to implement a simple ELF loader (you will extend the `main.c` file) and use it to load a simple ELF object file (the one compiled from `elf.c`). However, before starting on this lets make ourselves familiar with ELF files.

We provide a simple [Makefile](#) that compiles `elf.o` and `main` as ELF executables. Look over the makefile and then compile both files by running:

```
make
```

Lets take a look at the ELF files we compiled. We will use the `readelf` tool

```
$ readelf -a elf
```

ELF is the file format used for object files (`.o`'s), binaries, shared libraries and core dumps in Linux.

It's actually pretty simple and well thought-out.

ELF has the same layout for all architectures, however endianness and word size can differ; relocation types, symbol types and the like may have platform-specific values, and of course the contained code is arch specific.

The ELF files are used by two tools: 1) linker and 2) loader. A linker combines multiple ELF files into an executable or a library and a loader loads the executable ELF file in the memory of the process. On real operating systems loading may require relocation (e.g., if the file is dynamically linked it has to be linked again with all the shared libraries it depends on). In this homework we will not do any relocation (it's too complicated), we'll simply load an ELF file in memory and run it.

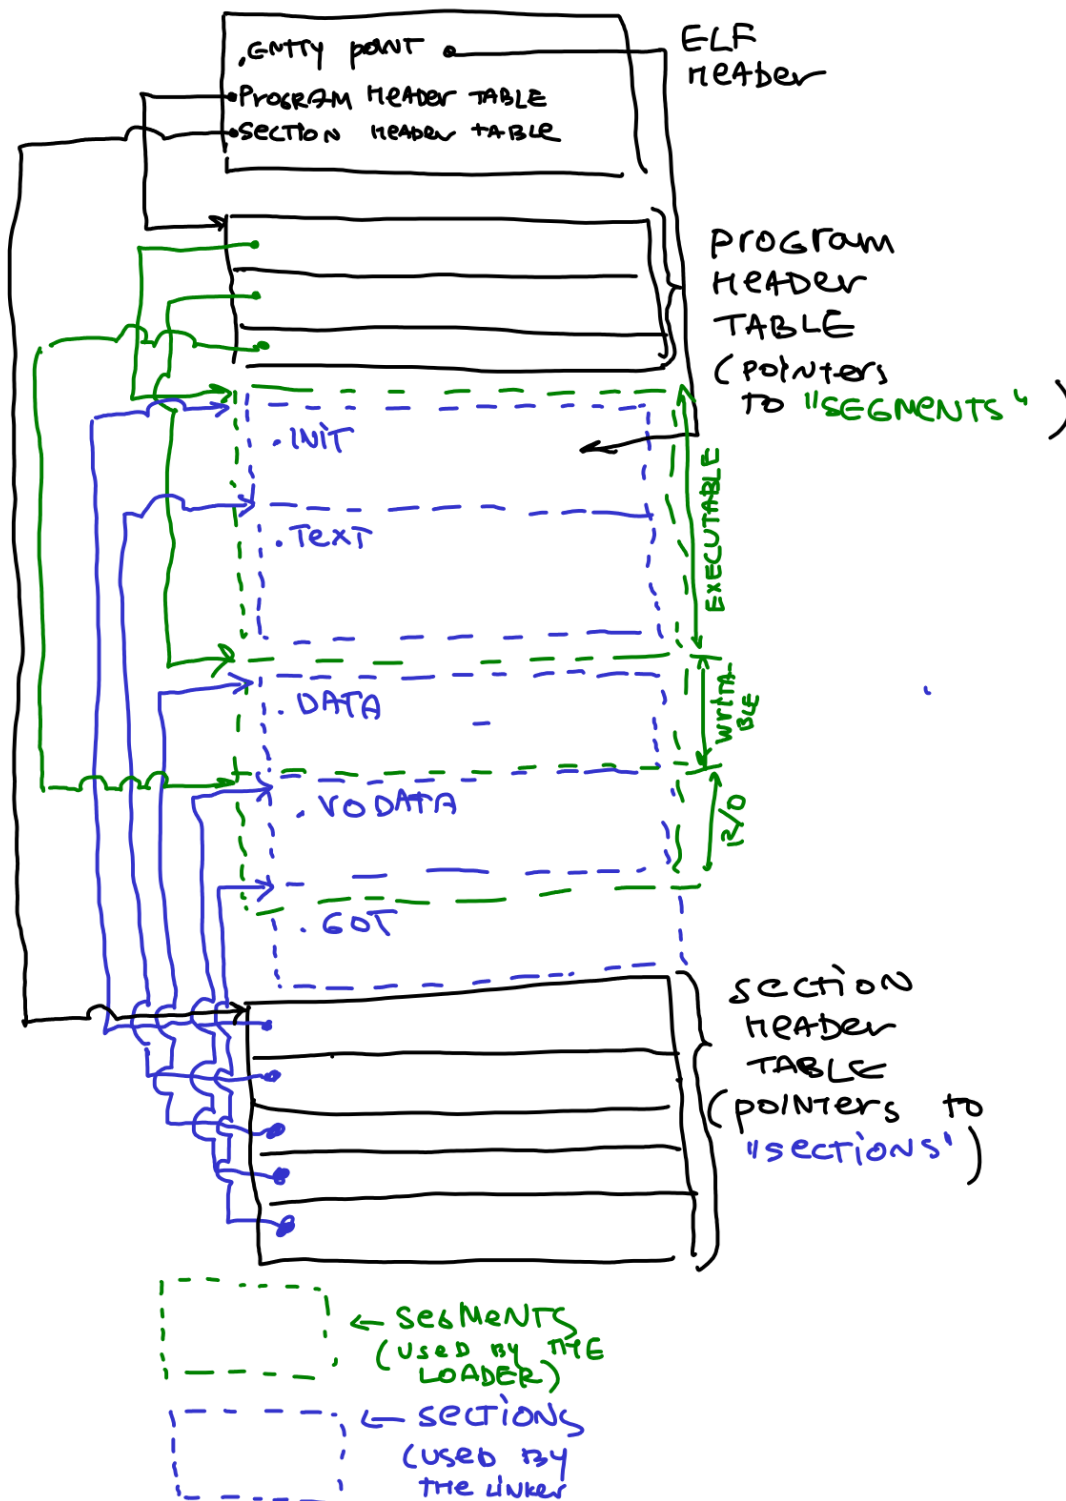
Linker and loader need two different views of the ELF file, i.e., they access it differently---the linker needs to know where the DATA, TEXT, BSS, and other sections are to merge them with sections

from other libraries. If relocation is required the linker needs to know where the symbol tables and relocation information is.

The loader, however, does not need any of these details. It simply needs to know which parts of the ELF file are code (executable), which are data and read-only data, and where to put the BSS in the memory of a process.

Hence the ELF file provides two separate views on the data inside the ELF file: 1) a more detailed view for the linker, and 2) a bit more high-level view for the loader. To provide these view each ELF file contains two arrays: Section Header Table (for the linker), and Program Header Table (for the loader). Both tables are simply arrays of entries that contain information about each part of the ELF file (e.g., where the sections for the linker and section for the loader are inside the ELF file).

Here is a simple figure of a typical ELF file that starts with the ELF header. The header contains pointers to the locations of Section Header Table and Program Header Table within the ELF file. Then each tables have entries that point to the starting locations of individual sections and segments.



Lets take a look at both arrays.

Linking view: Section Header Table (SHT)

The Section Header Table is an array in which every entry contains a pointer to one of the sections of the ELF file. It's a bit annoying but the parts of the ELF file used by the linker are called **"sections"**, and the parts used by the loader are called **"segments"** (my guess is that different CPU segments were configured in the past for each part of the program loaded in memory, hence

the name "segments", for example, an executable CPU segment was created for the executable parts of the ELF file (i.e., one segment that contained all executable sections like .text, .init, etc.).

Also don't get confused: sections and segments overlap. I.e., typically multiple sections (.text, .init) are all contained in one executable segment. Confusing, huh? It will become clear soon.

Lets take a look at what inside the ELF file. Run this command

```
readelf -a elf
```

If you scroll down to the **Section headers** you will see all "sections" of the ELF file that the linker can use:

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.text	PROGBITS	00000000	000074	00000d	00	WAX	0	0	1
[2]	.eh_frame	PROGBITS	00000010	000084	000038	00	A	0	0	4
[3]	.comment	PROGBITS	00000000	0000bc	00002b	01	MS	0	0	1
[4]	.debug_aranges	PROGBITS	00000000	0000e7	000020	00		0	0	1
[5]	.debug_info	PROGBITS	00000000	000107	000066	00		0	0	1
[6]	.debug_abbrev	PROGBITS	00000000	00016d	000055	00		0	0	1
[7]	.debug_line	PROGBITS	00000000	0001c2	000035	00		0	0	1
[8]	.debug_str	PROGBITS	00000000	0001f7	000106	01	MS	0	0	1
[9]	.symtab	SYMTAB	00000000	000300	0000e0	10		10	10	4
[10]	.strtab	STRTAB	00000000	0003e0	000024	00		0	0	1
[11]	.shstrtab	STRTAB	00000000	000404	000074	00		0	0	1

Key to Flags:

W (write), A (alloc), X (execute), M (merge), S (strings), I (info),
 L (link order), O (extra OS processing required), G (group), T (TLS),
 C (compressed), x (unknown), o (OS specific), E (exclude),
 p (processor specific)

Since `elf.c` is a very simple program, it has only .text section (i.e., code of the program), a bunch of sections that contain debugging information, and a .symtab section that contains imported and exported symbols.

Note, there is no .data or .bss sections for global variables (there are no globals in `elf.c`).

Moreover, since we linked `elf.c` to be a static executable, it is linked to run at address 0x0 (the `-Ttext 0` tells the linker to relocate the executable at linking time to work at 0x0):

```
elf: elf.o
      ld -m elf_i386 -N -e main -Ttext 0 -o elf elf.o

elf.o: elf.c
      $(CC) -c -fno-pic -static -fno-builtin -ggdb -m32 -fno-omit-frame-pointer elf.c
```

The symbol table contains these symbols

Symbol table '.symtab' contains 14 entries:

Num:	Value	Size	Type	Bind	Vis	Ndx	Name
0:	00000000	0	NOTYPE	LOCAL	DEFAULT	UND	
1:	00000000	0	SECTION	LOCAL	DEFAULT	1	
2:	00000010	0	SECTION	LOCAL	DEFAULT	2	
3:	00000000	0	SECTION	LOCAL	DEFAULT	3	
4:	00000000	0	SECTION	LOCAL	DEFAULT	4	
5:	00000000	0	SECTION	LOCAL	DEFAULT	5	
6:	00000000	0	SECTION	LOCAL	DEFAULT	6	
7:	00000000	0	SECTION	LOCAL	DEFAULT	7	
8:	00000000	0	SECTION	LOCAL	DEFAULT	8	
9:	00000000	0	FILE	LOCAL	DEFAULT	ABS	elf.c
10:	00000048	0	NOTYPE	GLOBAL	DEFAULT	2	__bss_start
11:	00000000	13	FUNC	GLOBAL	DEFAULT	1	main
12:	00000048	0	NOTYPE	GLOBAL	DEFAULT	2	_edata
13:	00000048	0	NOTYPE	GLOBAL	DEFAULT	2	_end

The `main` is our function (it's `FUNC`, and `GLOBAL`), the `__bss_start`, `_edata`, and `_end` are added by the linker to mark the start and end of the BSS, TEXT, and DATA sections.

If we take a look at the `main` executable, the ELF file is more complicated.

Section Headers:

[Nr]	Name	Type	Addr	Off	Size	ES	Flg	Lk	Inf	Al
[0]		NULL	00000000	000000	000000	00		0	0	0
[1]	.interp	PROGBITS	00000154	000154	000013	00	A	0	0	1
[2]	.note.ABI-tag	NOTE	00000168	000168	000020	00	A	0	0	4
[3]	.note.gnu.build-id	NOTE	00000188	000188	000024	00	A	0	0	4
[4]	.gnu.hash	GNU_HASH	000001ac	0001ac	000020	04	A	5	0	4
[5]	.dynsym	DYNSYM	000001cc	0001cc	000090	10	A	6	1	4
[6]	.dynstr	STRTAB	0000025c	00025c	0000b8	00	A	0	0	1
[7]	.gnu.version	VERSYM	00000314	000314	000012	02	A	5	0	2
[8]	.gnu.version_r	VERNEED	00000328	000328	000040	00	A	6	1	4
[9]	.rel.dyn	REL	00000368	000368	000058	08	A	5	0	4
[10]	.rel.plt	REL	000003c0	0003c0	000018	08	AI	5	22	4
[11]	.init	PROGBITS	000003d8	0003d8	000023	00	AX	0	0	4
[12]	.plt	PROGBITS	00000400	000400	000040	04	AX	0	0	16
[13]	.plt.got	PROGBITS	00000440	000440	000010	08	AX	0	0	8
[14]	.text	PROGBITS	00000450	000450	000222	00	AX	0	0	16
[15]	.fini	PROGBITS	00000674	000674	000014	00	AX	0	0	4
[16]	.rodata	PROGBITS	00000688	000688	000010	00	A	0	0	4
[17]	.eh_frame_hdr	PROGBITS	00000698	000698	000034	00	A	0	0	4
[18]	.eh_frame	PROGBITS	000006cc	0006cc	0000e0	00	A	0	0	4
[19]	.init_array	INIT_ARRAY	00001ecc	000ecc	000004	04	WA	0	0	4
[20]	.fini_array	FINI_ARRAY	00001ed0	000ed0	000004	04	WA	0	0	4
[21]	.dynamic	DYNAMIC	00001ed4	000ed4	000100	08	WA	6	0	4
[22]	.got	PROGBITS	00001fd4	000fd4	00002c	04	WA	0	0	4
[23]	.data	PROGBITS	00002000	001000	000008	00	WA	0	0	4
[24]	.bss	NOBITS	00002008	001008	000004	00	WA	0	0	1
[25]	.comment	PROGBITS	00000000	001008	00002b	01	MS	0	0	1
[26]	.debug_aranges	PROGBITS	00000000	001033	000020	00		0	0	1

[27]	.debug_info	PROGBITS	00000000	001053	000554	00	0	0	1	
[28]	.debug_abbrev	PROGBITS	00000000	0015a7	000132	00	0	0	1	
[29]	.debug_line	PROGBITS	00000000	0016d9	0000ee	00	0	0	1	
[30]	.debug_str	PROGBITS	00000000	0017c7	0003b4	01	MS	0	0	1
[31]	.symtab	SYMTAB	00000000	001b7c	000480	10	32	48	4	
[32]	.strtab	STRTAB	00000000	001ffc	00024f	00	0	0	1	
[33]	.shstrtab	STRTAB	00000000	00224b	00013c	00	0	0	1	

It contains all the section we've mentioned in class: `.text` (main code of the program), `.data` (data section for global variables), `.rodata` (data section for global read-only variables), `.bss` (uninitialized global variables), `.init` (init section to call the constructors that run before `main()`), `.got` (Global Offset Table), `.plt` (Procedure Linking Table for lazy linking of imported functions), and even the `.interp` (the section for the interpreter, i.e., the linker that links dynamically linked program before it runs, typically it's `/lib/ld-linux.so.2` on Linux systems).

Execution view: Program Header Table (PHT)

The Program Header Table contains information for the kernel on how to start the program. The `LOAD` directives determinate what parts of the ELF file get mapped into program memory.

Again, in our `elf` example the program header defines only two segments. And only one of them should be loaded by the operating system in memory to run.

Program Headers:

Type	Offset	VirtAddr	PhysAddr	FileSiz	MemSiz	Flg	Align
LOAD	0x000074	0x00000000	0x00000000	0x00048	0x00048	RWE	0x4
GNU_STACK	0x000000	0x00000000	0x00000000	0x00000	0x00000	RW	0x10

The only loadable section is linked to run at address `0x0`. We can inspect the `elf` binary with the `objdump` tool to see what is there:

```
$ objdump -d elf
```

```
elf:      file format elf32-i386
```

```
Disassembly of section .text:
```

```
00000000
```

```
:
```

```

0:  55                push    %ebp
1:  89 e5             mov     %esp, %ebp
3:  8b 55 08          mov     0x8(%ebp), %edx
6:  8b 45 0c          mov     0xc(%ebp), %eax
9:  01 d0             add     %edx, %eax
b:  5d                pop     %ebp
c:  c3                ret
```

Well, no surprises: it's the main function compiled into machine code.

Putting it all together: the ELF header

Neither the SHT nor the PHT have fixed positions, they can be located anywhere in an ELF file. To find them the ELF header is used, which is located at the very start of the file.

The first bytes contain the elf magic "\x7fELF", followed by the class ID (32 or 64 bit ELF file), the data format ID (little endian/big endian), the machine type, etc.

At the end of the ELF header are then pointers to the SHT and PHT. Specifically, the Section Header Table which is used by the linker starts at byte 1120 in the ELF file, and the Program Header Table starts at byte 52 (right after the ELF header)

```

ELF Header:
  Magic:      7f 45 4c 46 01 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF32
  Data:                               2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                               UNIX - System V
  ABI Version:                           0
  Type:                               EXEC (Executable file)
  Machine:                               Intel 80386
  Version:                               0x1
  Entry point address:                   0x0
  Start of program headers:               52 (bytes into file)
  Start of section headers:              1120 (bytes into file)
  Flags:                               0x0
  Size of this header:                    52 (bytes)
  Size of program headers:                32 (bytes)
  Number of program headers:               2
  Size of section headers:                40 (bytes)
  Number of section headers:              12
  Section header string table index:      11

```

Finally, the entry point of this file is at address `0x0`. This is exactly what we told the linker to do --- link the program to run at address `0x0`. And this is where the `main` function of the `elf.c` file is shown in the objdump.

Program loading in the kernel

The execution of a program starts inside the kernel, in the `exec("/bin/wc", ...)` system call takes a path to the executable file. The kernel reads the ELF header and the program header table (PHT), followed by lots of sanity checks.

The kernel then loads the parts specified in the `LOAD` directives in the PHT into memory. If an `INTERP` entry is present, the interpreter is loaded too. Statically linked binaries can do without an interpreter; dynamically linked programs always need `/lib/ld-linux.so` as interpreter because it includes some startup code, loads shared libraries needed by the binary, and performs relocations.

Finally control can be transferred to the entry point of the program or to the interpreter, if linking is required.

In case of a statically linked binary that's pretty much it, however with dynamically linked binaries a lot more magic has to go on.

First the dynamic linker (contained within the interpreter) looks at the `.dynamic` section, whose address is stored in the PHT.

There it finds the `NEEDED` entries determining which libraries have to be loaded before the program can be run, the `*REL*` entries giving the address of the relocation tables, the `VER*` entries which contain symbol versioning information, etc.

So the dynamic linker loads the needed libraries and performs relocations (either directly at program startup or later, as soon as the relocated symbol is needed, depending on the relocation type).

Finally control is transferred to the address given by the symbol `_start` in the binary. Normally some gcc/glibc startup code lives there, which in the end calls `main()`.

What you need to do in this homework: load an ELF file

While ELF might look a bit intimidating, in practice the loading algorithm is trivial:

1. Read the ELF header (This [Wiki](#) page should help).
2. One of the ELF header fields tells you the offset of the program header table inside the file.
3. Read each entry of the program header table (i.e., read each program header)
4. Each program header has an offset and size of a specific segment inside the ELF file (e.g., a executable code). You have to read it from the file and load it in memory.
5. When done with all segments, jump to the entry point of the program. (Note since we don't control layout of the address space at the moment, we load the sections at some random place in memory (the place that is allocated for us by the `mmap()` function). Obviously the address of the entry point should be an offset within that random area. Such loading will not work for a real ELF file, but ours is simple: it's statically linked, and contains the code that can run at any location in memory. So even though it's linked to run at `0x0` it will run anywhere where you load it.

Looks manageable. We make a couple of simplifications. First we create a very simple ELF file out of `elf.c` --- it contains only one function, it has no data, and the code can be placed anywhere in memory and will run ok (it simply does not refer to any global addresses --- all variables are on the stack).

The `main.c` file provides definitions for the structs that match the ELF header and entries of the Program Header Table. So you can simply read the header out of the ELF file by using `open`, `lseek(fd, offset, SEEK_SET)` to set the position of where to read to be `offset`, and `read(fd, &elf, sizeof(elfhdr))`, and then you can read the offset of the program header table, get the number of the entries in the program header table from the ELF header, and read all entries of the program header table one by one using `lseek`, and `read` in a similar fashion as before. If the entry has `ELF_PROG_LOAD` type, you will load it in memory.

To load the segment in memory, you can allocate executable memory with the following function


```
code_va = mmap(NULL, ph.memsz, PROT_READ | PROT_WRITE | PROT_EXEC,
               MAP_ANONYMOUS | MAP_PRIVATE, 0, 0);
```

where `ph.memsz` is the size of the segment that we currently are loading.

You should figure out where the entry point of your program is (it is in one of the segments, and since you map the segments at the addresses returned by the `mmap()` function, you need to find where it ends up in your address space.

If you found the entry point then type cast it to the function pointer that matches the function signature of the `sum` function and call it.

```
if (entry != NULL) {
    sum = entry;
    ret = sum(1, 2);
    printf("sum:%d\n", ret);
};
```

Your program should take the name of the elf file as first argument and return the following result:

```
$ ./main elf
sum:3
```

Extra credit: (10% bonus)

Try loading [elf-data.c](#) file. Can you explain why it crashes?

Fix the crash by performing a relocation step for the ELF file when it is loaded.

These [ELF tutorial](#) and [Executable and Linkable Format 101 Part 3: Relocations](#) resources can be helpful.

Submit your work

Submit your solution through Gradescope [Gradescope CS143A Operating Systems](#). Please zip all of your files (`main.c`, `Makefile`) and submit them. If you have done extra credit then place `main.c` and `Makefile` for extra credit part into folder "extra" The structure of the zip file should be the following:

```
/
- Makefile
- main.c
- /extra --optional
  - Makefile --optional
  - main.c --optional
```